Product ⌄   Use cases ⌄   Resources ⌄   Company ⌄   Enterprise   Pricing

Clay University > AI Research & Writing
> AI Writing at Scale: Best Practices

AI Writing at Scale: Best Practices

Author   Yash Tekriwal

▶

💬 Prompting ≠ Prompting *at scale*

There's a million resources on learning prompt engineering, but *almost none* on prompt engineering *at scale*. When you're writing one prompt to try and run as accurately as possible for hundreds, or

three best practices for you to use to make your prompts as efficient as possible.

1️⃣ Sequence

When you're prompting, giving steps in logical, precise order increases clarity and reduces the rate of hallucination (made-up results). By giving AI an order to how you want it to execute, you're giving it guardrails on how to "think", instead of letting it run wild.

Here's an example of an un-sequenced prompt:

- Write me an email first line to ⎡Name⎤ using the profile information below:
  ┌─────────────────────────┐
  │ Enrich Person from   │
  │ Profile   │
  └─────────────────────────┘

Here's a prompt with some better sequencing:

- I will give you the full profile for ⎡Name⎤ below. Write a first line for an email intro that provides context on why I'm reaching out, personalized to the prospect. In order of priority, mention awards, volunteering

Search ⌘ + K

🎓 All Categories

🏪 Find Companies

🏭 Enrich Companies

👥 Find People

🧑‍💼 Enrich People

✨ AI Research & Writing

POPULAR LESSONS

AI Writing at Scale: Best Practices

Enrich & Validate Emails - Waterfall

Finding Companies

Pushing data to your CRM

🌈   Product ⌄   Use cases ⌄   Resources ⌄   Company ⌄   Enterprise   Pricing

Enrich Person from
Profile

2️⃣ Chunk

Chunking is the most counterintuitive principle to most users. We are primed to think that *more is better*, but as with many things in life, *less is more* when it comes to prompting at scale.

Many people think that the more context and instructions you give to AI, the better it will perform.

Not necessarily.

Working with AI is a lot like training a new employee. If you give a new employee a 20-page guide on how to write cold emails on Day 1, you're not setting them up for success. The same is true with AI.

With onboarding new employees, the goal is to be as insights rich and compact as possible. Again, the same is true with AI.

This means removing filler words and fluff, but it goes even further.

communicate the exact
same meaning in fewer
words, always use the
option with fewer
words.

That goes as far as using I will
provide instead of I am going to give
you. Remember that LLMs aren't
true "intelligence". They're complex
algorithms trained on billions of data
points, and using that massive
dataset to predict the next most
likely words from your prompt. So
the less room there is for confusion -
the better.

3️⃣ Sample

Whenever you want a narrow range
of results, it's helpful to provide a
narrow range of examples as
inspiration.

Think about onboarding a new
employee again. If you give them a
set of principles, but then don't give
them any examples of what great
looks like - there's still a lot of room
for error.

examples of what great looks like - they have something to reference when thinking about quality.

Once again, AI is exactly the same. Provides examples of the kinds of output you want it to create makes the range of results much more narrow, and tightens the guardrails of what great can look like.

4️⃣ Specify

Last of the prompting principles, but potentially most important, is *specify*.

When you're onboarding a new employee, principles and a mission are good enough to get going, but if you don't also provide a set of rules to *not* violate, you can't blame them for trying.

You have to give people a set of "Dos" and "Do Nots" to set them up for maximum success. AI operates in exactly the same way.

When you specify guardrails like "Never use 1st person pronouns" or "Always start with Just read, came

final output.

You should always consider specifying:

- Objective
- Tone
- Length of response
- Format of response

5️⃣ Iterate

The final principle is actually the one to rule them all. Iterate iterate iterate.

Nobody nails the perfect prompt down on the first try. Write your best first draft, run it on 10 rows of data, and look at what you don't like. Consider the edge cases. Add specificity, chunks, examples, or sequencing to your prompt to improve it. Repeat.

Do this over and over again until you've got 10 rows of data returning exactly what you want. Then, run it for the next 10. Repeat.

Once you've got a prompt reliable running for 10-20 rows of data, the

Product ⌄   Use cases ⌄   Resources ⌄   Company ⌄   Enterprise   Pricing

How was this lesson?

Let us know if there is anything we can do to improve the design, structure, format, you name it!

Share feedback ↗

More on AI Research & Writing

Product ⌄   Use cases ⌄   Resources ⌄   Company ⌄   Enterprise   Pricing

▶ Overview: AI Research & Writing

▶ ChatGPT vs. Claygent

▶ AI Writing at Scale: Best Practices | You're Here |

outbound motion in seconds, not months

14 day free Pro trial - No credit card required

Try Clay free →

Product	Blog	Tools	Support	Company
Login	AI for Sales Prospecting	Email Finder	Send us an email	Customer Stories
University	Why good CRM data matters	Headcount Finder	Join Slack	Wall of Love
Integrations		Credits Calculator	FAQ	Careers
Pricing	Google Maps lead generation in 5 easy steps	Chrome extension	Status	Creators Program
Changelog				Integrate with Clay

Automation

AI Lead
Generation

CPM Calculator

Margin
Calculator

Headcount
Directory

Glossary

Dossier

Job Board

Legal

Privacy Policy

Terms of Service

Do not sell my data

clay

© Clay 2024 / Uplevel your data enrichment. Scale personalized outreach.